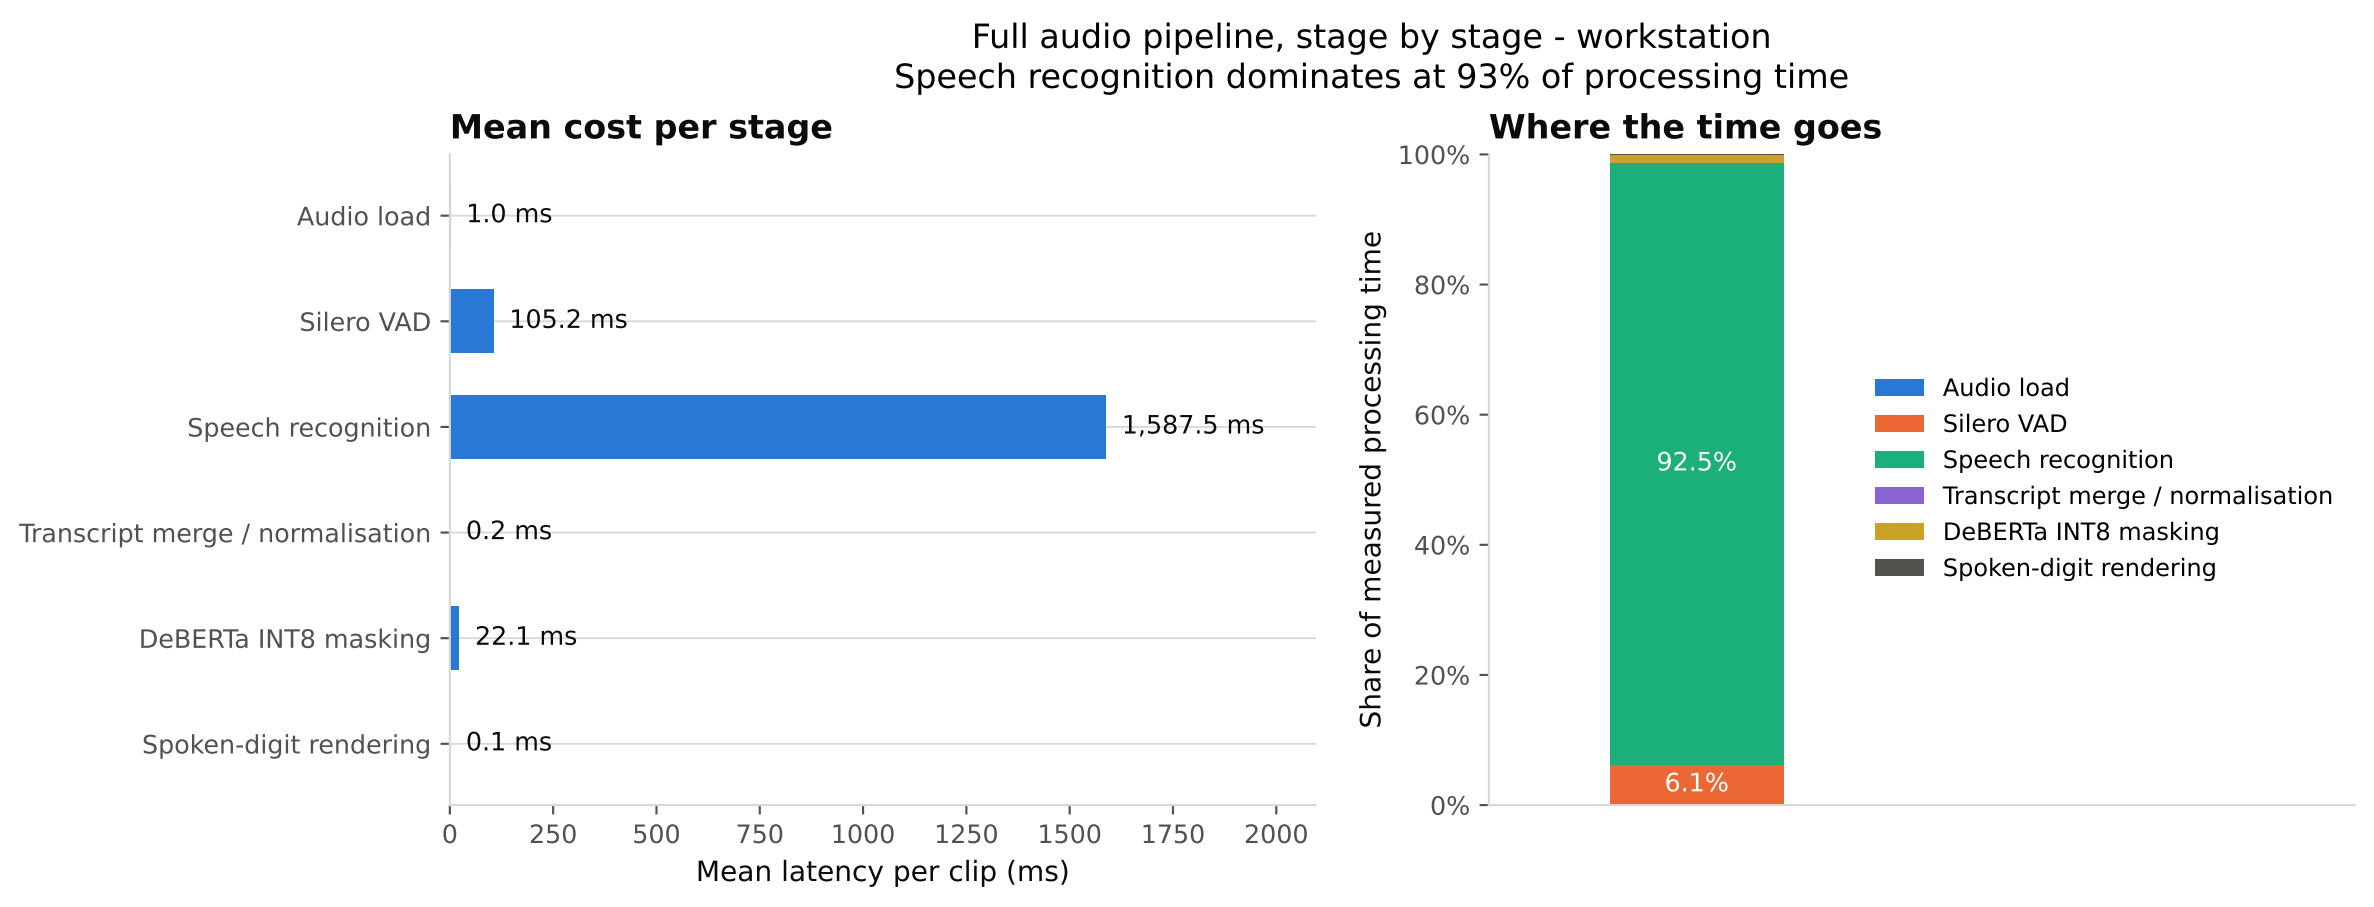
Full audio pipeline, stage by stage - workstation
Speech recognition dominates at 93% of processing time
Mean cost per stage
Audio load
Silero VAD
Speech recognition
Transcript merge / normalisation
DeBERTa INT8 masking
Spoken-digit rendering
1.0 ms
105.2 ms
1,587.5 ms
0.2 ms
22.1 ms
0.1 ms
0
250
500
750
1000
1250
1500
1750
2000
Mean latency per clip (ms)
Where the time goes
Share of measured processing time
100%
80%
60%
40%
20%
0%
92.5%
6.1%
Audio load
Silero VAD
Speech recognition
Transcript merge / normalisation
DeBERTa INT8 masking
Spoken-digit rendering