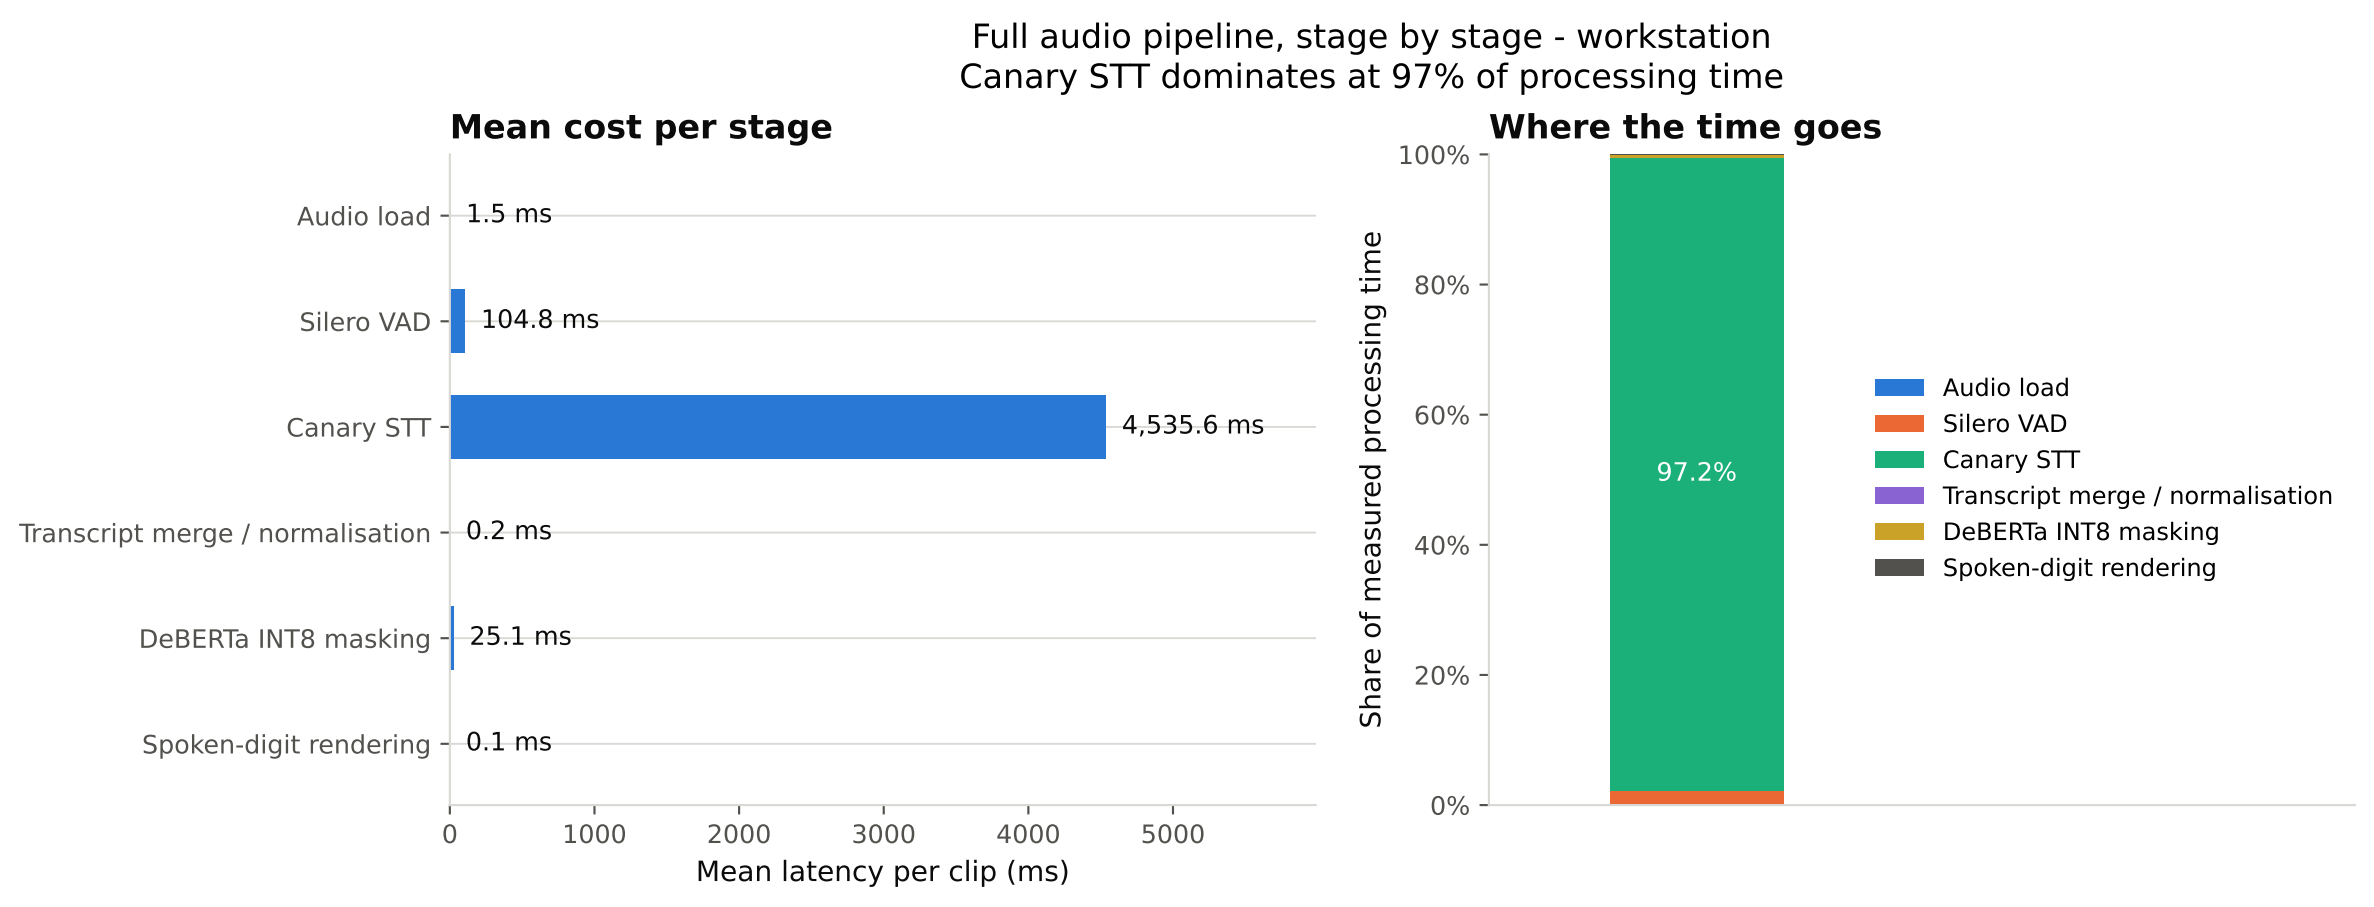
Full audio pipeline, stage by stage - workstation
Canary STT dominates at 97% of processing time
Mean cost per stage
Audio load
Silero VAD
Canary STT
Transcript merge / normalisation
DeBERTa INT8 masking
Spoken-digit rendering
1.5 ms
104.8 ms
4,535.6 ms
0.2 ms
25.1 ms
0.1 ms
0
1000
2000
3000
4000
5000
Mean latency per clip (ms)
Where the time goes
Share of measured processing time
100%
80%
60%
40%
20%
0%
97.2%
Audio load
Silero VAD
Canary STT
Transcript merge / normalisation
DeBERTa INT8 masking
Spoken-digit rendering