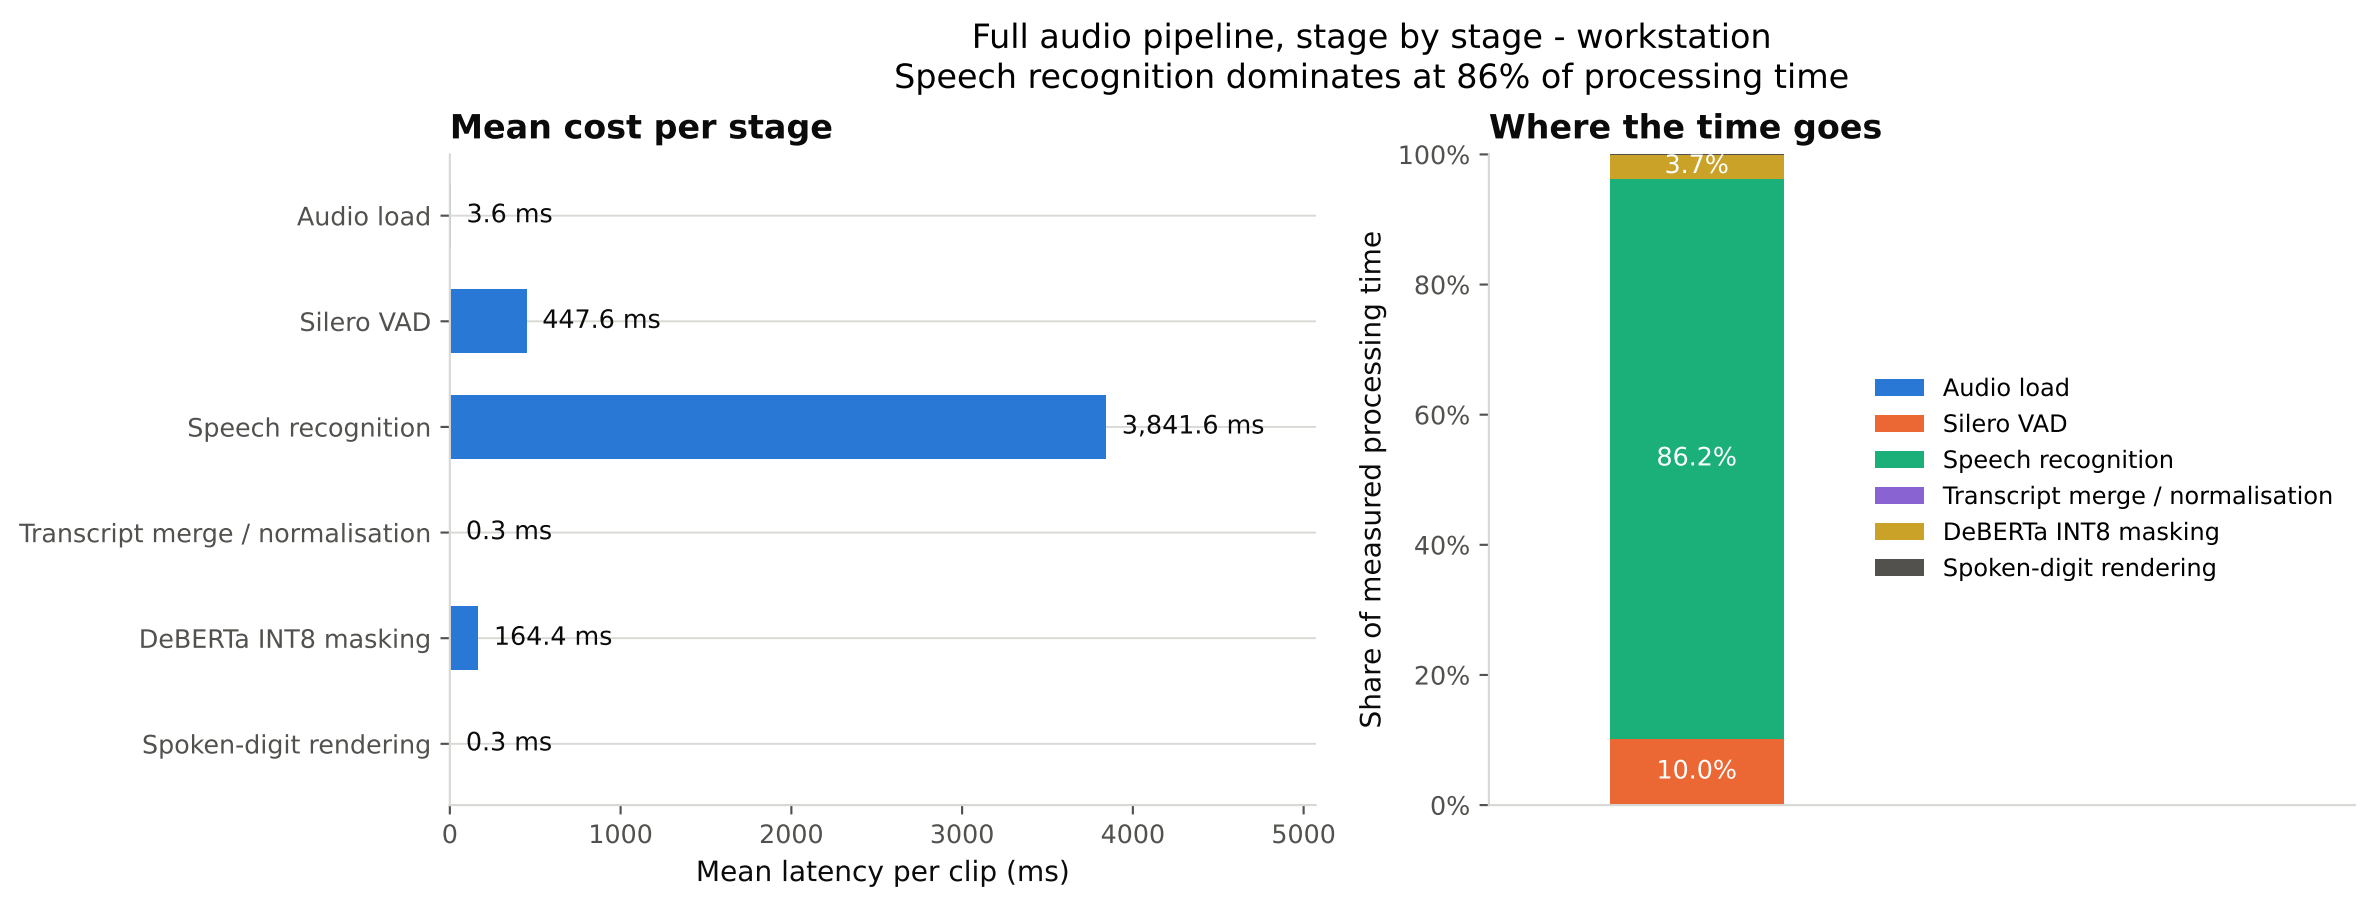
Full audio pipeline, stage by stage - workstation
Speech recognition dominates at 86% of processing time
Mean cost per stage
Audio load
Silero VAD
Speech recognition
Transcript merge / normalisation
DeBERTa INT8 masking
Spoken-digit rendering
3.6 ms
447.6 ms
3,841.6 ms
0.3 ms
164.4 ms
0.3 ms
0
1000
2000
3000
4000
5000
Mean latency per clip (ms)
Where the time goes
Share of measured processing time
100%
80%
60%
40%
20%
0%
3.7%
86.2%
10.0%
Audio load
Silero VAD
Speech recognition
Transcript merge / normalisation
DeBERTa INT8 masking
Spoken-digit rendering