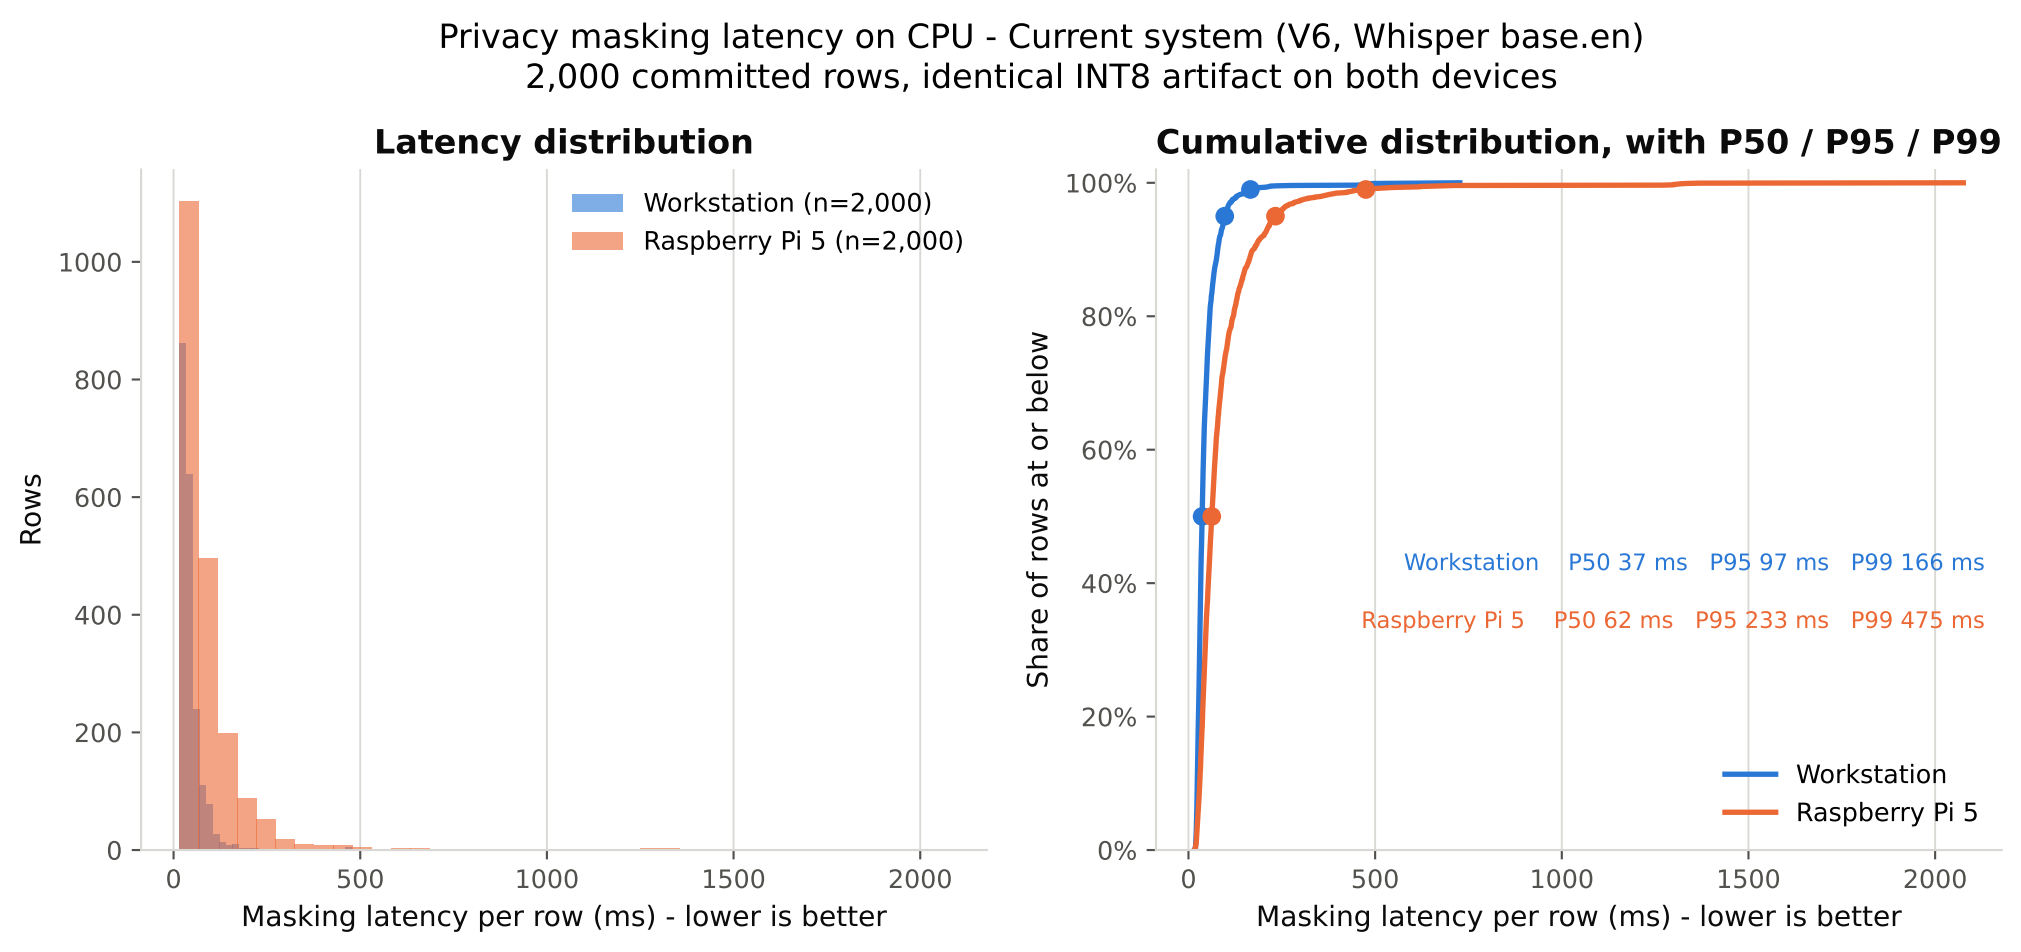
Privacy masking latency on CPU - Current system (V6, Whisper base.en)
2,000 committed rows, identical INT8 artifact on both devices
Latency distribution
Workstation (n=2,000)
Raspberry Pi 5 (n=2,000)
Rows
0
200
400
600
800
1000
0
500
1000
1500
2000
Masking latency per row (ms) - lower is better
Cumulative distribution, with P50 / P95 / P99
Share of rows at or below
0%
20%
40%
60%
80%
100%
Workstation P50 37 ms P95 97 ms P99 166 ms
Raspberry Pi 5 P50 62 ms P95 233 ms P99 475 ms
Workstation
Raspberry Pi 5
0
500
1000
1500
2000
Masking latency per row (ms) - lower is better