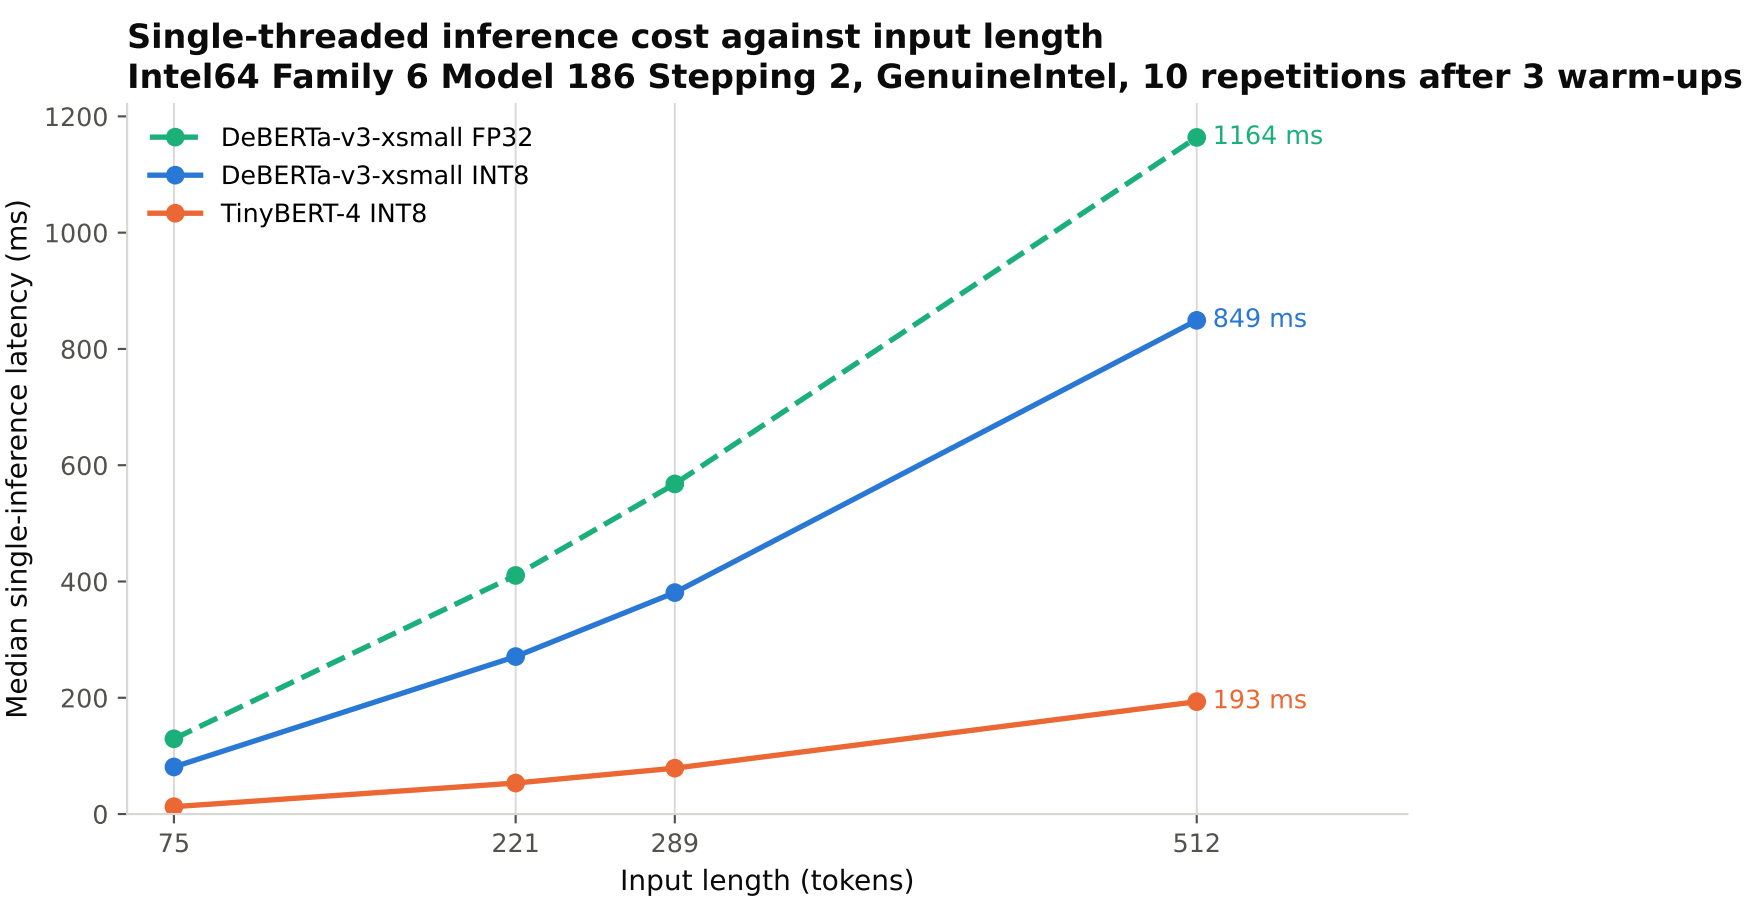
Single-threaded inference cost against input length
Intel64 Family 6 Model 186 Stepping 2, GenuineIntel, 10 repetitions after 3 warm-ups
DeBERTa-v3-xsmall FP32
DeBERTa-v3-xsmall INT8
TinyBERT-4 INT8
1164 ms
849 ms
193 ms
Median single-inference latency (ms)
0
200
400
600
800
1000
1200
75
221
289
512
Input length (tokens)