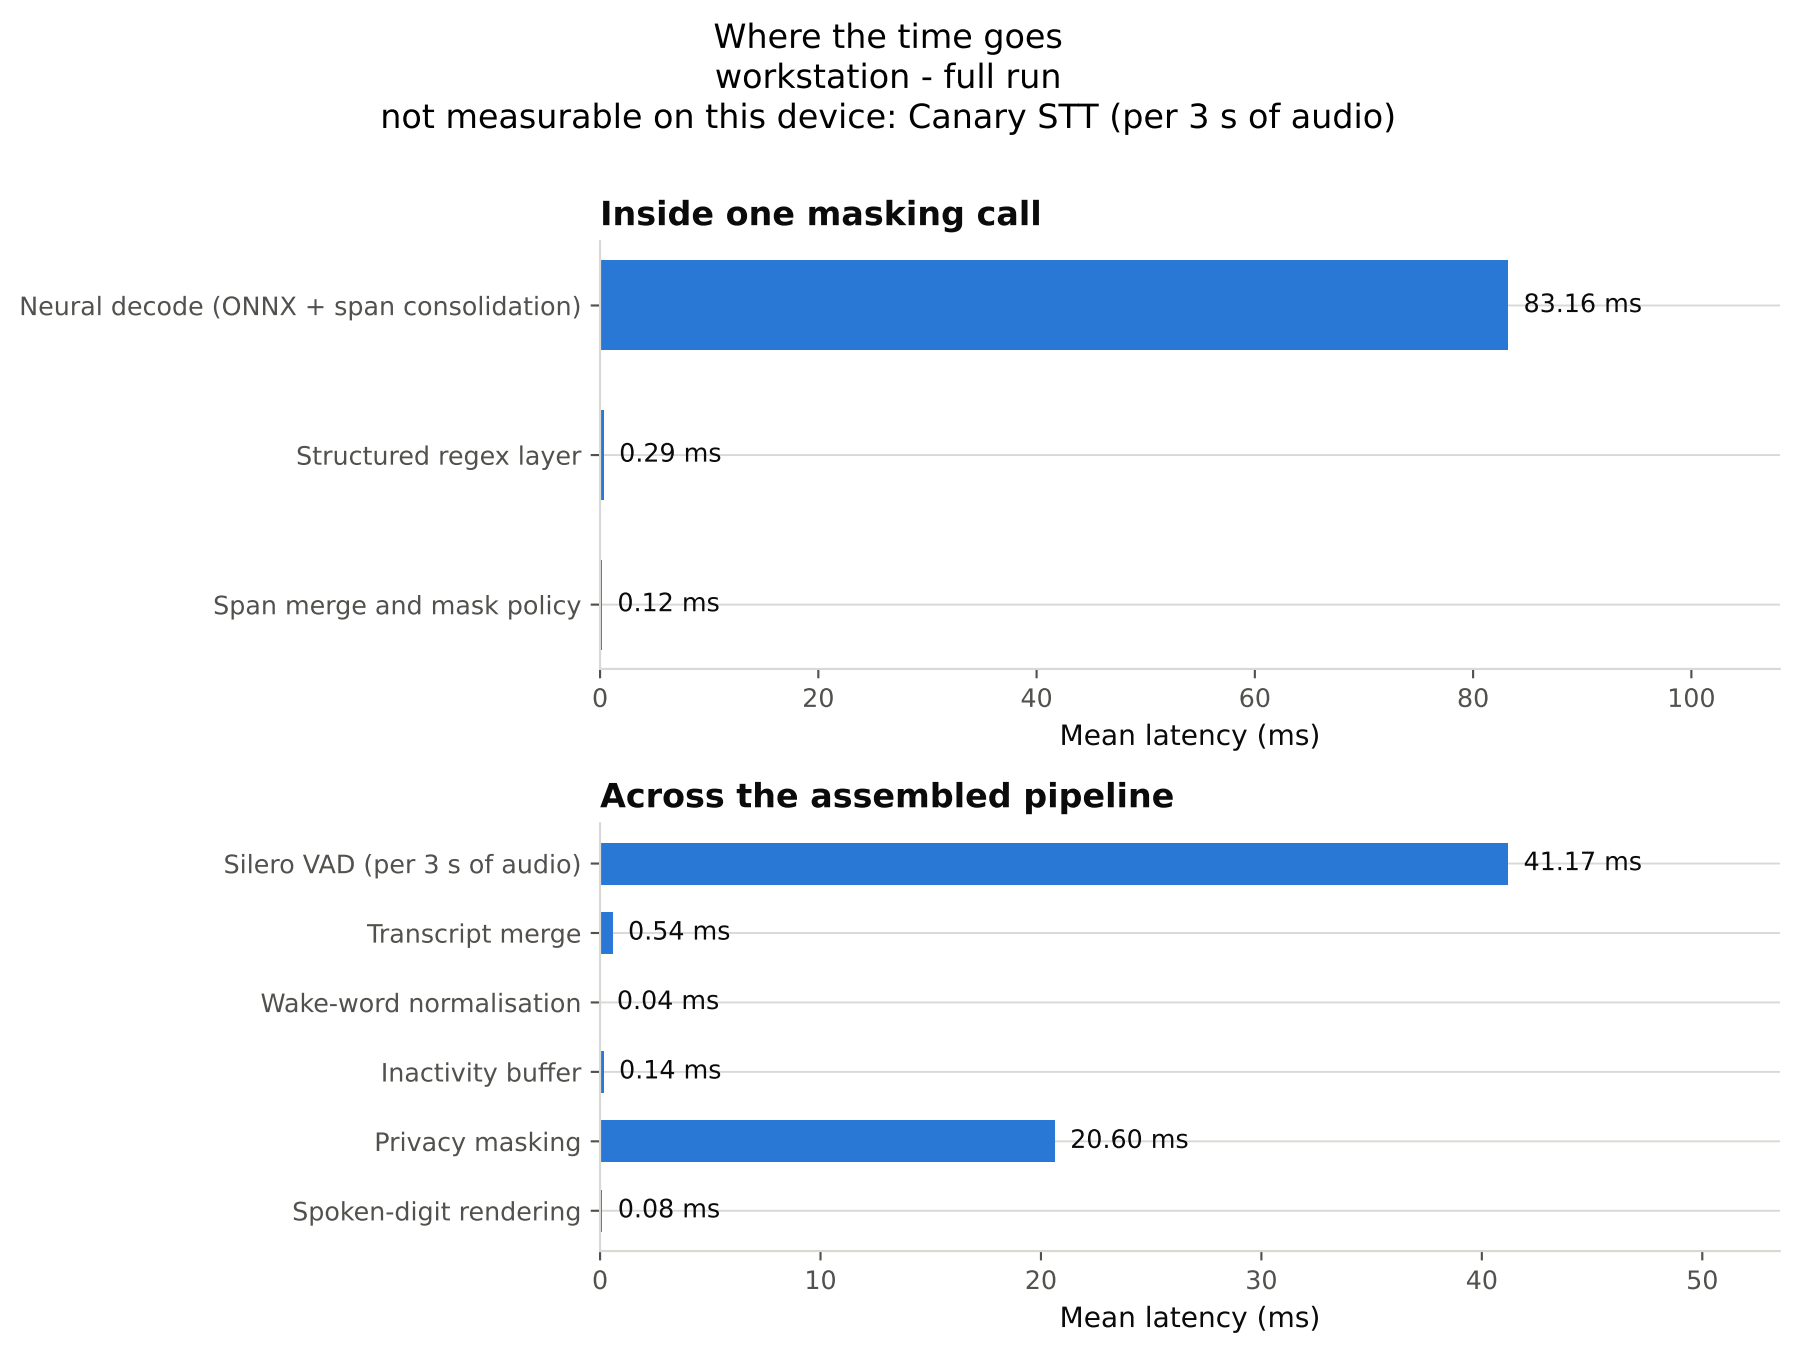
Where the time goes
workstation - full run
not measurable on this device: Canary STT (per 3 s of audio)
Inside one masking call
Neural decode (ONNX + span consolidation)
83.16 ms
Structured regex layer
0.29 ms
Span merge and mask policy
0.12 ms
0
20
40
60
80
100
Mean latency (ms)
Across the assembled pipeline
Silero VAD (per 3 s of audio)
41.17 ms
Transcript merge
0.54 ms
Wake-word normalisation
0.04 ms
Inactivity buffer
0.14 ms
Privacy masking
20.60 ms
Spoken-digit rendering
0.08 ms
0
10
20
30
40
50
Mean latency (ms)